Real Estate Sales Prediction Analysis

Chuqi(Angel) Jin

Contents

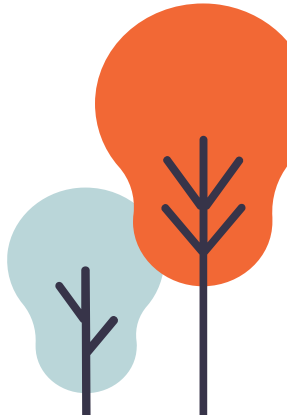

Introduction

Dataset: Real Estate Sales (Kaggle)

Goals: Which factors significantly impact property sale price?

- **Hypothesis:** The year of the real estate transaction has a significant on the sale price.

- **Original Dataset:** The dataset initially included real estate transaction data from Connecticut between 2009 till 2022, with a total of 10,000 observations and 12 variables.

- **Cleaned Dataset:** After carefully cleaning and selecting, I got 2574 observations and 12 variables.
  - Remove all 0 and NA values.
  - Filter the observations to the years 2017–2022.
  - Divided the sale price and estimated values by 1,000.

- **Model Building Dataset**: each categorical variables has too many category, subset top 2 of each categorical variables and left dataset with 397 variables for model building

Data Description

Response Variable

Sale.Price: (measured in thousand dollars) – The actual sale price of the property

- **Date**: The property transaction date (Numerical)
- **Year**: The property transaction year (Numerical)
- **Locality**: The property locality/area (Categorical)
- **Estimated.Value**: (measured in thousand dollars) The estimated value of the property (Numerical)
- **Property**: Types of properties suitable for various family sizes (Categorical)
- **Residential**: Indicates whether the property is designated for residential use (Categorical)
- **num_rooms**: The number of rooms in the property (Numerical)
- **num_bathrooms**: The number of bathrooms in the property (Numerical)
- **carpet_area**: (measured in square feet) The carpet area of the property (Numerical)
- **Property_tax_rate**: Tax rate applied to the property's assessed value (Numerical)
- **Face**: Direction the main entrance/facade of the property is oriented towards (Categorical)

EDA: Response Variable

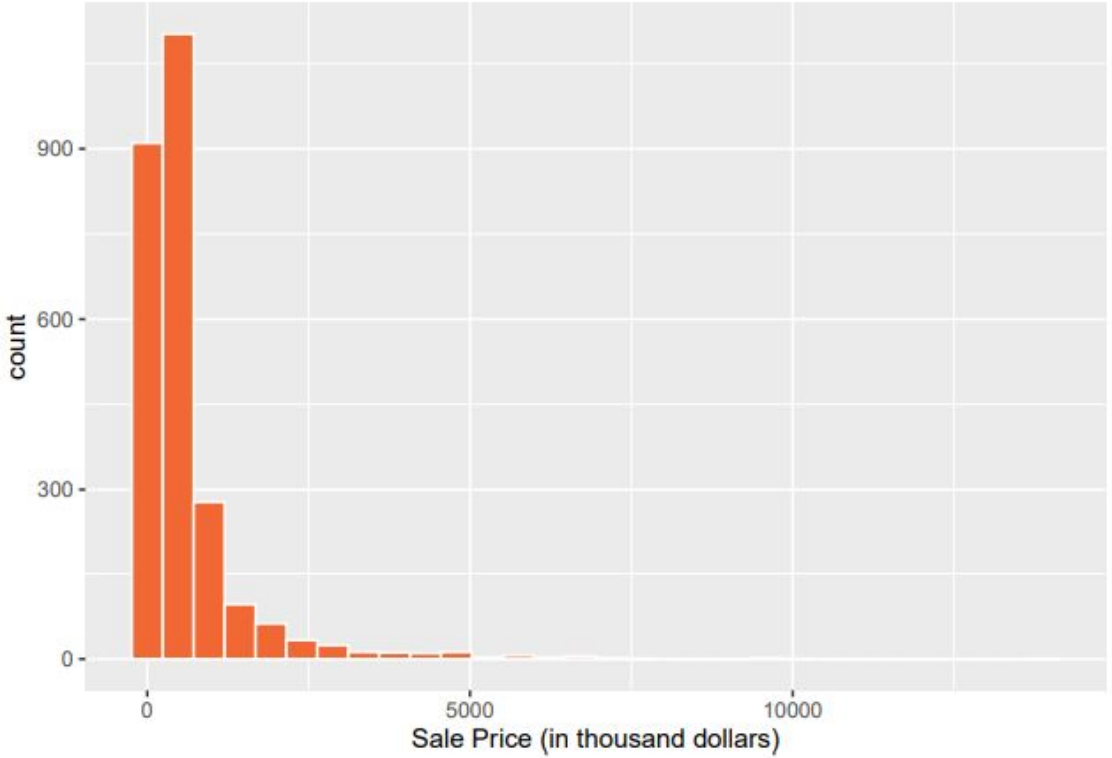

Sale.Price

- The histogram indicates that the response variable is not normally distributed.
- Right skewed

EDA: Predictor Variables

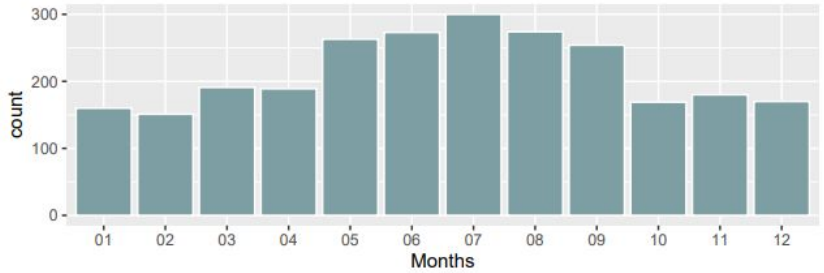

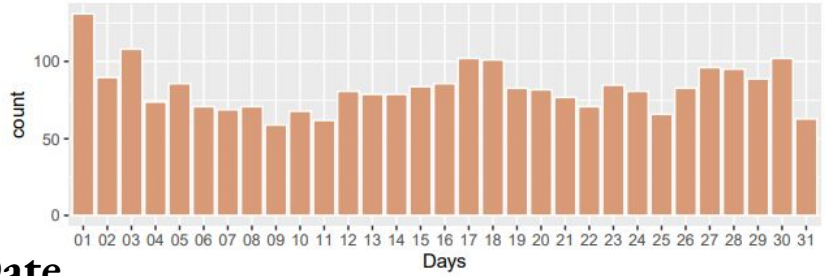

Date

– July is the most popular month for real estate sales.
– The most real estate sales occurred on the first day of the month.

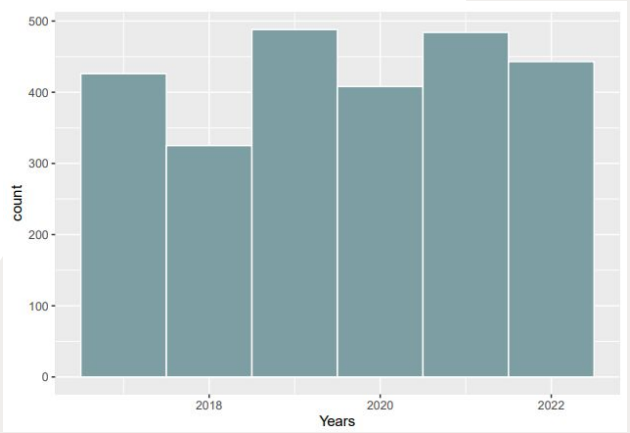

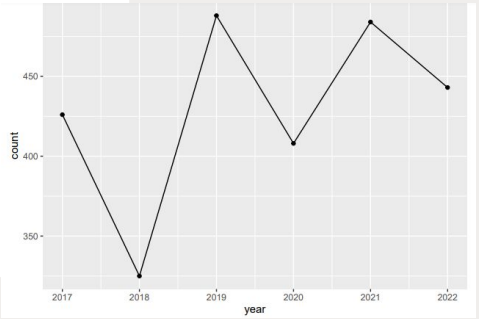

Year

– Real estate sales were highest in 2019 and lowest in 2018.

EDA

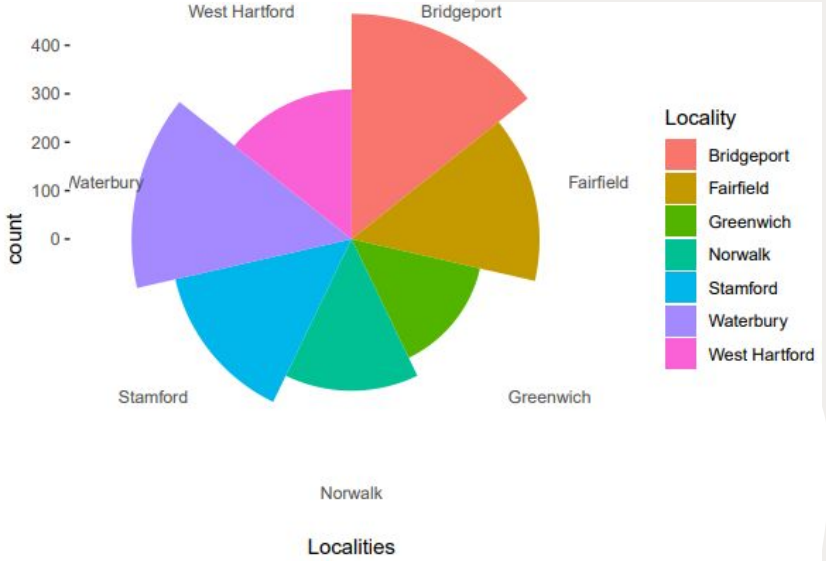

Locality

– 6 different locations

– The top 2 localities are Bridgeport & Waterbury

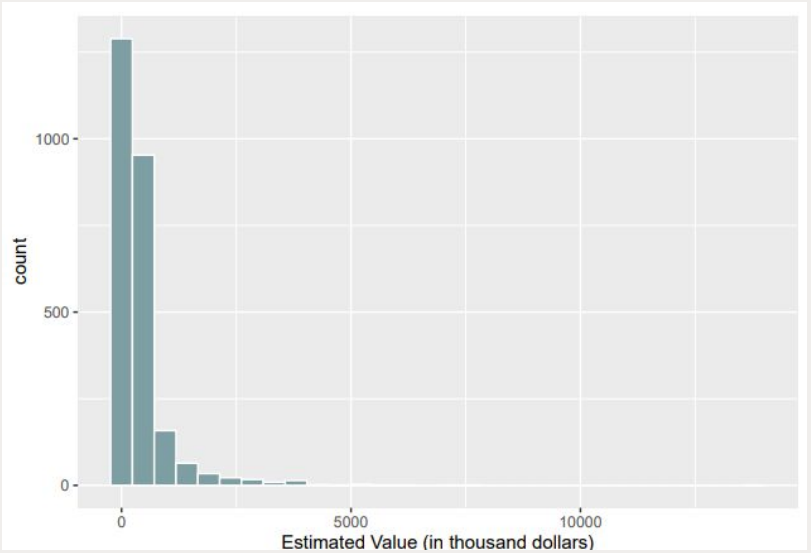

Estimated Value

- The histogram indicates that this variable is not normally distributed.
- Right skewed

EDA

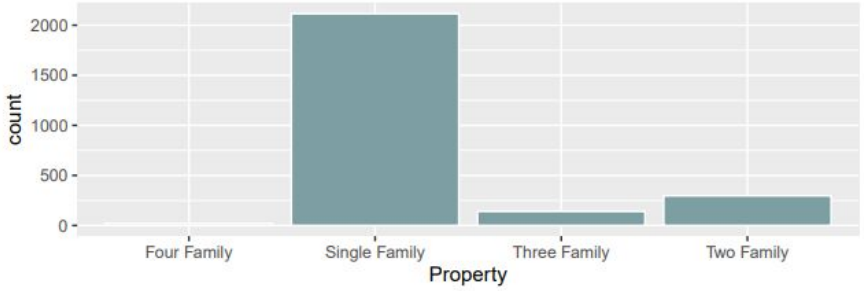

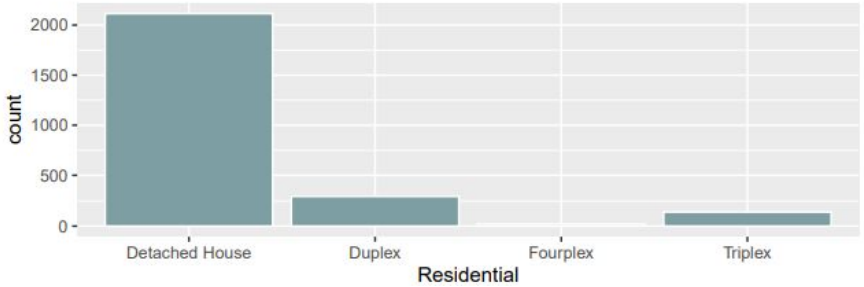

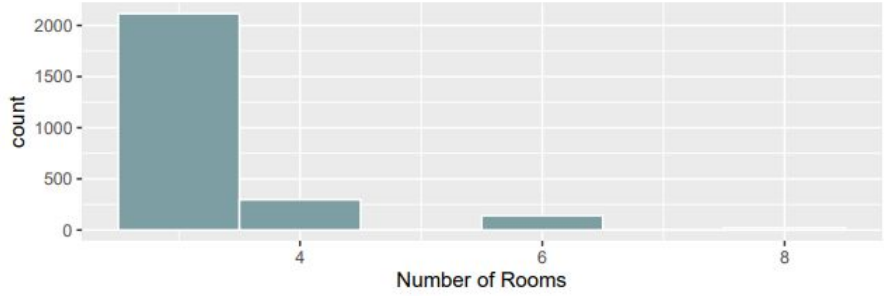

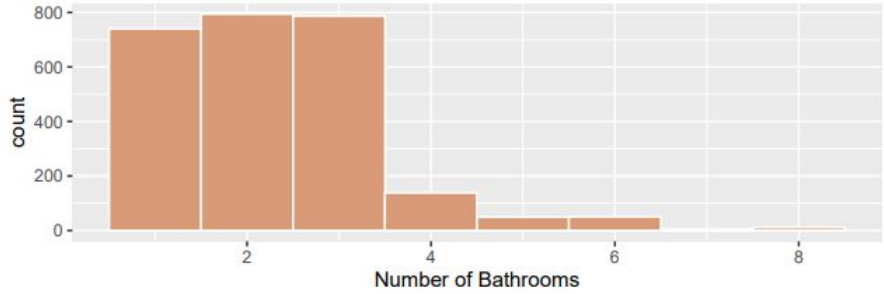

Property
– 4 different property types
– The top 2 properties are Single & Two Family

Residential
– 4 different residential types
– The top 2 residence are Detached House & Duplex

Room numbers
– Between 3 and 8 rooms

Bathroom numbers
– Between 1-8 bathrooms

EDA

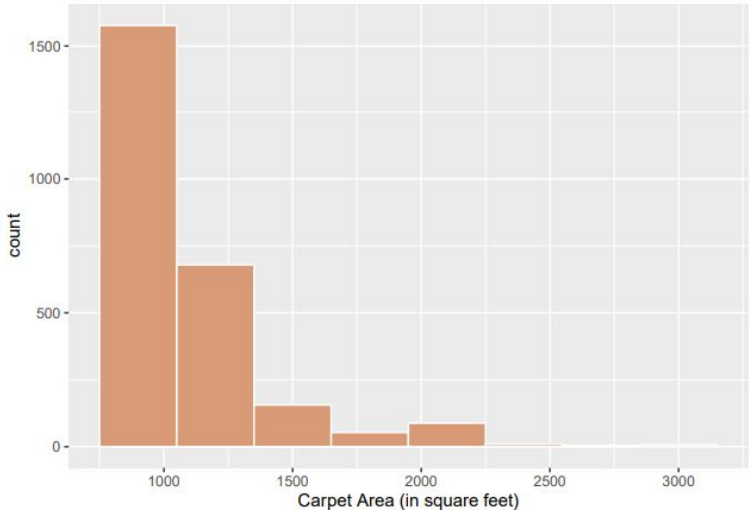

Carpet Area

– Range of 900 - 2989 square feets

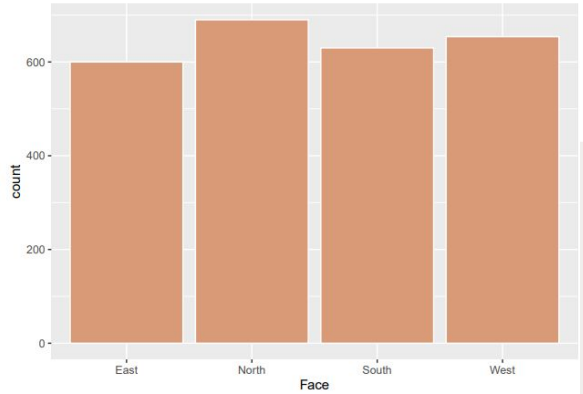

Face

– 4 different facing directions

– The top 2 directions are North & West

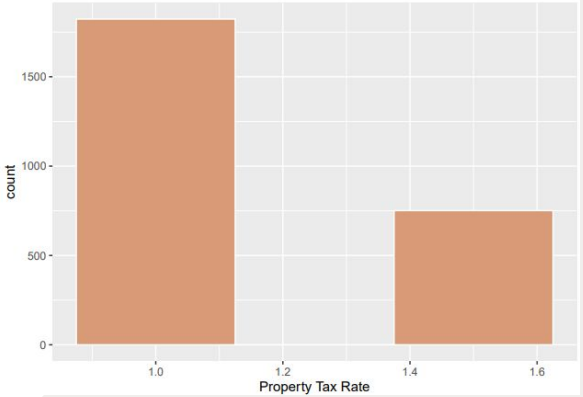

Property Tax Rate

– Range from 1.004 to 1.422

EDA

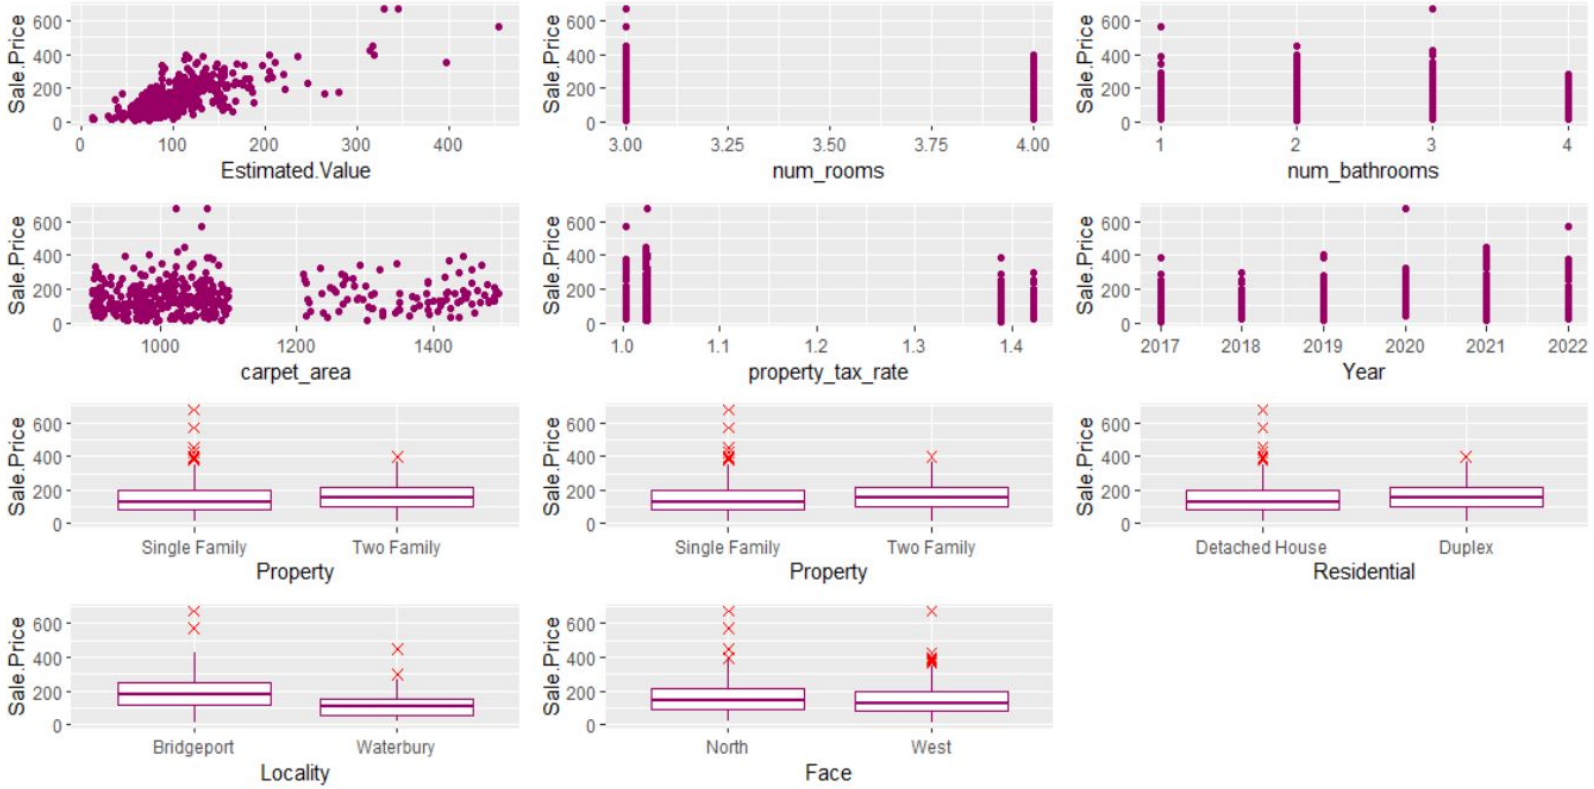
Sale.Price
Estimated.Value
num_rooms
num_bathrooms
carpet_area
property_tax_rate
Year
Property
Single Family
Two Family
Residential
Detached House
Duplex
Locality
Bridgeport
Waterbury
Face
North
West

Multicollinearity Analysis

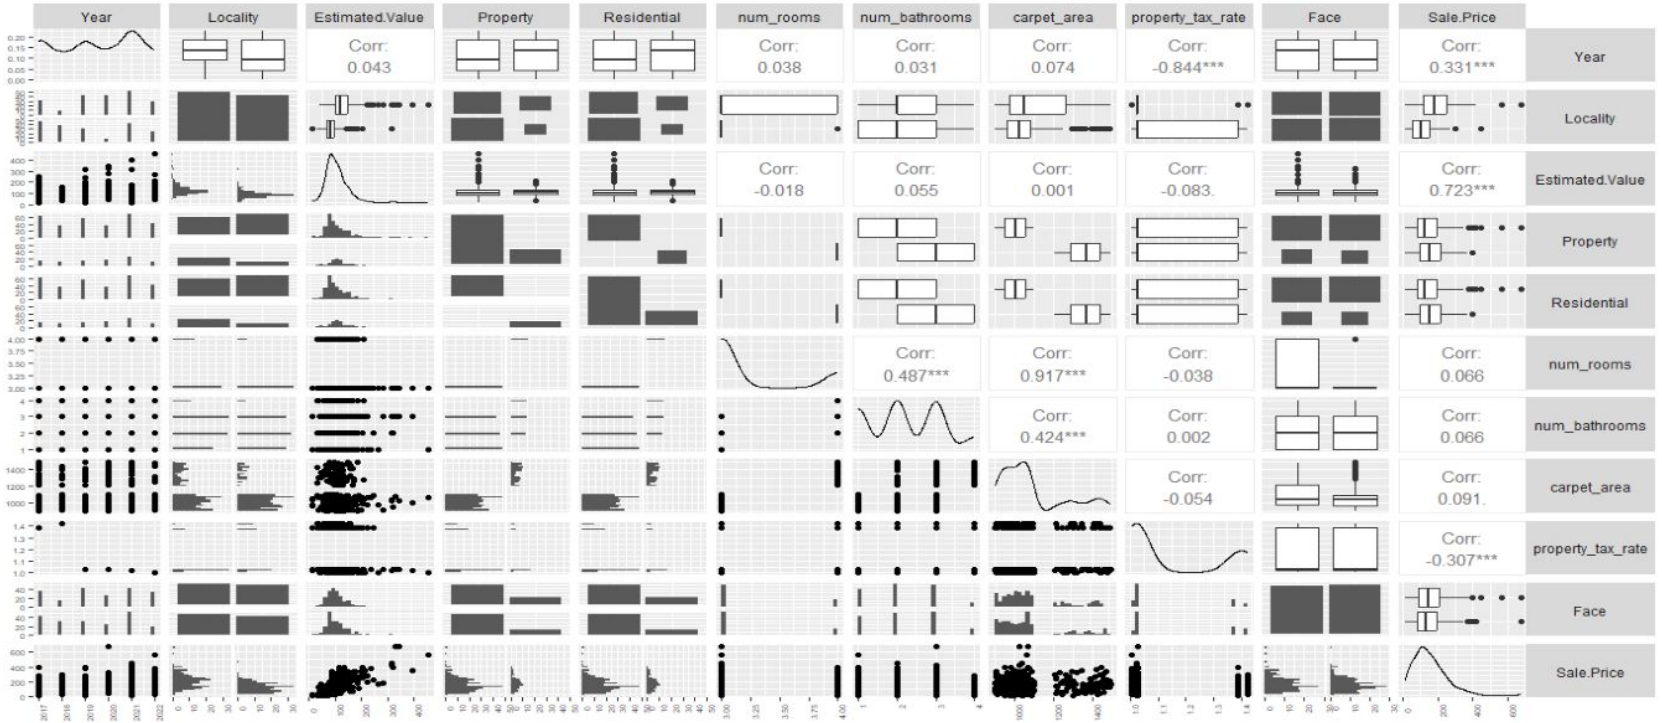

VIF

```
## Coefficients: (2 not defined because of singularities)
```

Year	Locality	Estimated.Value	num_rooms
3.581429	1.416801	1.268067	7.034652
num_bathrooms	carpet_area	property_tax_rate	Face
1.337689	6.421536	3.757161	1.021831

Remove Property & Residential

Remove num_rooms

Year	Locality	Estimated.Value	num_bathrooms
3.549281	1.390207	1.248429	1.231240
carpet_area	property_tax_rate	Face	
1.250881	3.749222	1.016746	

Forward Stepwise Model

```
Call:
lm(formula = Sale.Price ~ Estimated.Value + Year + Locality +
    carpet_area, data = modelbuilding)

Residuals:
     Min       1Q   Median       3Q      Max
-223.530  -36.063   -0.146   29.334  238.186

Coefficients:
                   Estimate Std. Error t value Pr(>|t|)
(Intercept)       -3.220e+04  3.587e+03  -8.977   <2e-16 ***
Estimated.Value    1.263e+00  6.457e-02  19.558   <2e-16 ***
Year               1.594e+01  1.776e+00   8.974   <2e-16 ***
LocalityWaterbury -1.516e+01  6.866e+00  -2.207   0.0279 *
carpet_area        3.324e-02  1.833e-02   1.813   0.0705 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 59.77 on 392 degrees of freedom
Multiple R-squared:  0.6226, Adjusted R-squared:  0.6187
F-statistic: 161.7 on 4 and 392 DF,  p-value: < 2.2e-16
```

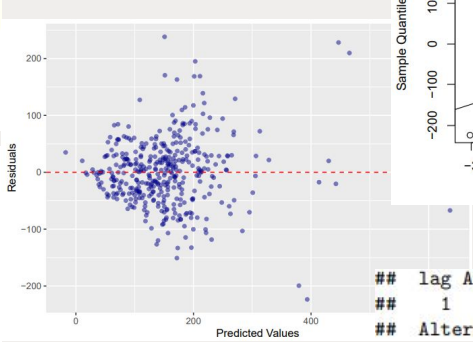

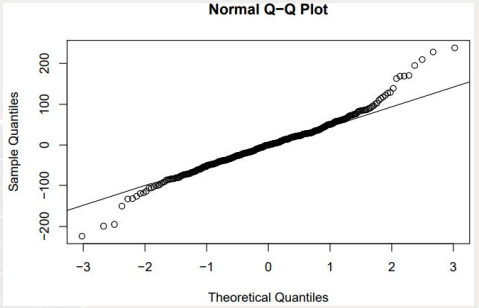


```
##  lag Autocorrelation D-W Statistic p-value
##    1      -0.04604595       2.086949    0.39
##  Alternative hypothesis: rho != 0
```

- Final predictors: Estimated value, year, locality and carpet area.
- Adjusted R^2: 0.6187
- Not pass Normality, Linearity or Constant Variance
- Pass independence check

Interaction Term

- Convert all categorical variables into binary variables
- Check interaction between Estimated Value & Locality
- Check the effect of Year
    - Present but does not seem to significantly impact the model
- Decides to transform Sale.Price (y) using log transformation
- Recheck the interaction between Estimated Value & Locality
- Recheck the effect of Year
- Decide to remove carpet area, number of bathrooms and facing direction from the predictor variables

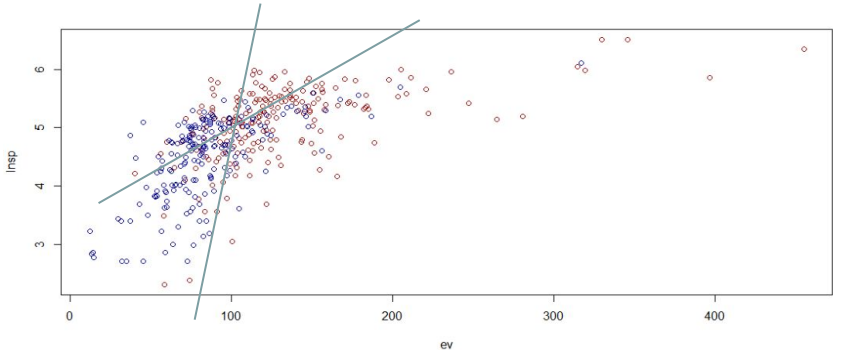

Model 1

ln(Sale Price) = Year + Locality + Estimated Value + Property Tax Rate

```
Call:
lm(formula = lnsp ~ 0 + year + ev + L + ptr)

Residuals:
    Min      1Q  Median      3Q     Max
-2.3570 -0.2865  0.1087  0.3693  1.1493

Coefficients:
       Estimate Std. Error t value Pr(>|t|)
year  2.572e-03  9.319e-05  27.601  < 2e-16 ***
ev    7.708e-03  5.763e-04  13.375  < 2e-16 ***
L    -1.536e-01  6.197e-02  -2.478   0.0136 *
ptr  -1.009e+00  1.573e-01  -6.414 4.07e-10 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 0.5348 on 393 degrees of freedom
Multiple R-squared:  0.988,     Adjusted R-squared:  0.9879
F-statistic:  8099 on 4 and 393 DF,  p-value: < 2.2e-16
```

- Adjusted R-squared: 0.9879
- Residuals not normally distributed
- Residuals satisfied independence test
- Residuals satisfied constant variance test

```
        Shapiro-Wilk normality test

data:  m1$residuals
W = 0.94203, p-value = 2.47e-11
```

```
lag Autocorrelation D-W Statistic p-value
  1       0.0873208      1.815225   0.054
Alternative hypothesis: rho != 0
```

```
data:  m1
BP = 2.4551, df = 3, p-value = 0.4835
```

Model 2

ln(Sale Price) = Year + Locality * Estimated Value + Property Tax Rate

- Adjusted R-squared: 0.9886
- Residuals not normally distributed
- Residuals failed independence test, so it shows autocorrelation
- Residuals satisfied constant variance test

```
Call:
lm(formula = lnsp ~ 0 + year + ev * L + ptr)

Residuals:
     Min       1Q   Median       3Q      Max
-2.45139 -0.27082  0.09784  0.35141  1.14163

Coefficients:
       Estimate Std. Error t value Pr(>|t|)
year  2.660e-03  9.194e-05  28.936  < 2e-16 ***
ev    5.986e-03  6.514e-04   9.189  < 2e-16 ***
L    -7.913e-01  1.379e-01  -5.736 1.94e-08 ***
ptr  -9.663e-01  1.527e-01  -6.328 6.79e-10 ***
ev:L  6.516e-03  1.269e-03   5.135 4.45e-07 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.5183 on 392 degrees of freedom
Multiple R-squared:  0.9888, Adjusted R-squared:  0.9886
F-statistic:  6903 on 5 and 392 DF,  p-value: < 2.2e-16
```

```
##  Shapiro-Wilk normality test
##
## data:  lm4$residuals
## W = 0.94502, p-value = 5.706e-11
```

```
##  studentized Breusch-Pagan test
##
## data:  lm4
## BP = 4.1493, df = 4, p-value = 0.3862
```

```
##  lag Autocorrelation D-W Statistic p-value
##    1       0.1078089       1.773688   0.022
##  Alternative hypothesis: rho != 0
```

Model Comparison & Model Diagnostics

```
Analysis of Variance Table

Model 1: lnsp ~ 0 + year + ev + L + ptr
Model 2: lnsp ~ 0 + year + ev * L + ptr
  Res.Df    RSS Df Sum of Sq      F    Pr(>F)
1    393 112.40
2    392 105.31  1    7.0848 26.371 4.448e-07 ***
```

- P-value is extremely small, therefore, model 1 vs. model 2 indicates that there is a huge difference between 2 models.
- Choose Model 2 as our final model

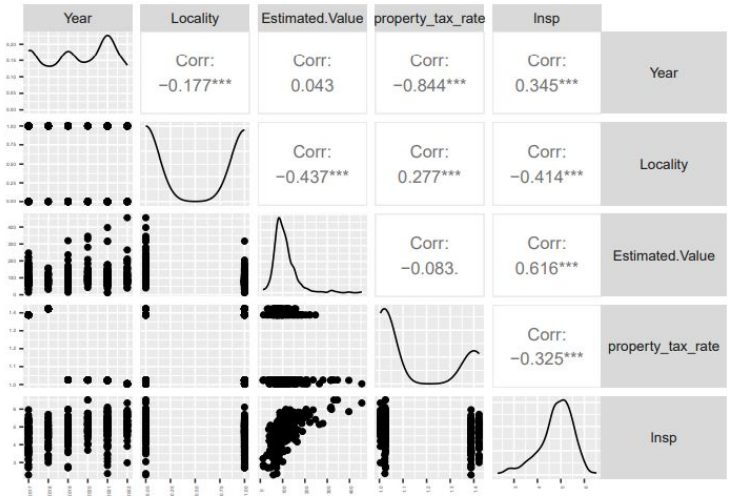


```
    year       ev        L      ptr    ev:L
50.939679 9.072707 13.670066 45.890357 9.769706
```

Final Model

$$\ln(\text{Sale Price}) = 2.660e^{-3}\text{Year} + 5.986e^{-3}\text{ Estimated Value} - 7.913e^{-1}\text{Locality} + 6.516e^{-3}(\text{Locality} * \text{ Estimated Value}) - 9.6636e^{-1}\text{Property Tax Rate}$$

```
Call:
lm(formula = lnsp ~ 0 + year + ev * L + ptr)

Residuals:
     Min       1Q   Median       3Q      Max
-2.45139 -0.27082  0.09784  0.35141  1.14163

Coefficients:
      Estimate Std. Error t value Pr(>|t|)
year  2.660e-03  9.194e-05  28.936  < 2e-16 ***
ev    5.986e-03  6.514e-04   9.189  < 2e-16 ***
L    -7.913e-01  1.379e-01  -5.736 1.94e-08 ***
ptr  -9.663e-01  1.527e-01  -6.328 6.79e-10 ***
ev:L  6.516e-03  1.269e-03   5.135 4.45e-07 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.5183 on 392 degrees of freedom
Multiple R-squared:  0.9888, Adjusted R-squared:  0.9886
F-statistic:  6903 on 5 and 392 DF,  p-value: < 2.2e-16
```

Conclusion & Limitation

Conclusion: The final model I obtained has a very high adjusted R-squared. The final model suggests that a real estate sales price is influenced by the year, estimated value, locality, property tax rate, and the interaction of estimated value and location. Based on my hypothesis, I can conclude that the year of the real estate transaction has a significant effect on the sale price.

Limitation: My final model does not satisfy the normality or independence assumptions, and it has a high vif score with several variables, which could be due to my dataset being too small. As a result, for my future work, I'd like to collect additional observations, expand the dataset size, and rebuild the model to see if I can achieve better results for residual assumption checks.

THANK YOU!

Questions?